Lokesh Reddy Maddali

Data Engineer | Azure Databricks & PySpark | Lakehouse & ETL/ELT | AWS | RAG & LLM Pipelines

+91 9346320355 | lokeshreddymaddali@gmail.com | Bangalore | LinkedIn | GitHub

CORE SKILLS

- Databricks & PySpark Pipelines
- ETL/ELT & Data Integration
- Delta Lake & Lakehouse Design
- Spark Performance Tuning
- Data Modeling (Kimball, Data Vault)
- Data Quality & Governance
- Production Support & RCA
- CI/CD & Infrastructure as Code
- RAG & LLM Data Pipelines

TECHNICAL SKILLS

Languages: Python, SQL, PySpark, Scala

Azure: Databricks, Data Factory, ADLS Gen2, Synapse, Azure SQL, Event Hub, Key Vault, Fabric, Purview

Databricks: Delta Lake, Medallion architecture, Structured Streaming, Asset Bundles, workspace administration

AWS: S3, Glue, Lambda, Step Functions, Redshift, Kinesis, EMR

Data Platforms: Snowflake, BigQuery, PostgreSQL, SQL Server, Oracle

Orchestration: Airflow, dbt, Control-M

DevOps: Git, Azure DevOps, GitHub Actions, Terraform, Docker

Governance: Purview, Collibra, data lineage, role-based access control

AI: RAG, embeddings, pgvector, LangChain, LLM evaluation

CERTIFICATIONS

- Microsoft Certified: Fabric Data Engineer Associate (DP-700)
- AWS Certified Data Engineer – Associate
- Microsoft Certified: Azure Data Fundamentals (DP-900)
- Microsoft Certification: DP-750

EDUCATION

Master's in Information Technology
Webster University, USA | 2024 – 2025

B.E., Electronics & Communication
R.M.D. Engineering College, Chennai | 2021

PROFILE SUMMARY

Data Engineer with 3+ years of professional experience and a 14-month data engineering internship in the USA. I build and run production pipelines on Azure Databricks and PySpark, design Lakehouse platforms on Delta Lake, and keep critical data reliable through tuning, monitoring, and root cause analysis. Certified in Microsoft Fabric (DP-700) and AWS Data Engineering, with hands-on experience building RAG and LLM data pipelines.

WORK EXPERIENCE

Data Engineer | Cenit Technologies | Oct 2025 – Present | ***Bangalore, India***

- Build and maintain **10+** production pipelines on Azure Databricks and Data Factory, ingesting **10+** sources and datasets over **1 TB**.
- Designed a Lakehouse on Delta Lake and ADLS using a Medallion (bronze/silver/gold) model, with dimensional and Data Vault 2.0 layers.
- Tuned Spark jobs (joins, caching, data skew), cutting pipeline run time by **~30%** and cloud cost by **~15%**.
- Added validation, reconciliation, and automated remediation, reducing pipeline failures by **~30%**; provide L2/L3 support with root cause analysis.
- Implemented governance with Purview, lineage, and role-based access control, and automated deployments with Azure DevOps, Asset Bundles, and Terraform.
- Automated **~3 hours/week** of manual work, working closely with architects and business stakeholders on requirements.

Data Engineer | Infosys | Jan 2022 – Dec 2023 | ***Chennai, India***

- Designed and maintained **10+** ETL pipelines integrating multiple source systems, including legacy systems and stored procedures.
- Cut ETL processing time by **35%** through SQL optimization and workflow refactoring.
- Delivered the data behind the AVON automated dashboard, used by **10+** stakeholders.

INTERNSHIP

Data Engineering Intern | Travelers Insurance | Aug 2024 – Sep 2025 | ***USA***

- Built Python and SQL ETL pipelines for insurance policy and claims data across **5+** datasets, with validation and reconciliation for accurate reporting.

PROJECTS

India Air Quality Pipeline | Python, PostgreSQL, GitHub Actions, Streamlit | GitHub

- Built an hourly, serverless pipeline tracking **20** Indian cities (bronze, silver, gold), computing the official CPCB AQI from 24-hour PM averages with idempotent upserts.
- Added automated data quality gates, z-score spike alerts, and forecast-accuracy tracking; batched API calls cut run time from **282s to 32s**.

DocQA: RAG Platform | Python, FastAPI, PostgreSQL/pgvector, Groq LLM | GitHub

- Built an idempotent RAG pipeline (chunking, embeddings, HNSW index) over **~80** Wikipedia articles with an API returning cited LLM answers.
- Evaluated retrieval on a **98**-question LLM-generated test set: hit@1 of **97.96%** and MRR of **0.98**.